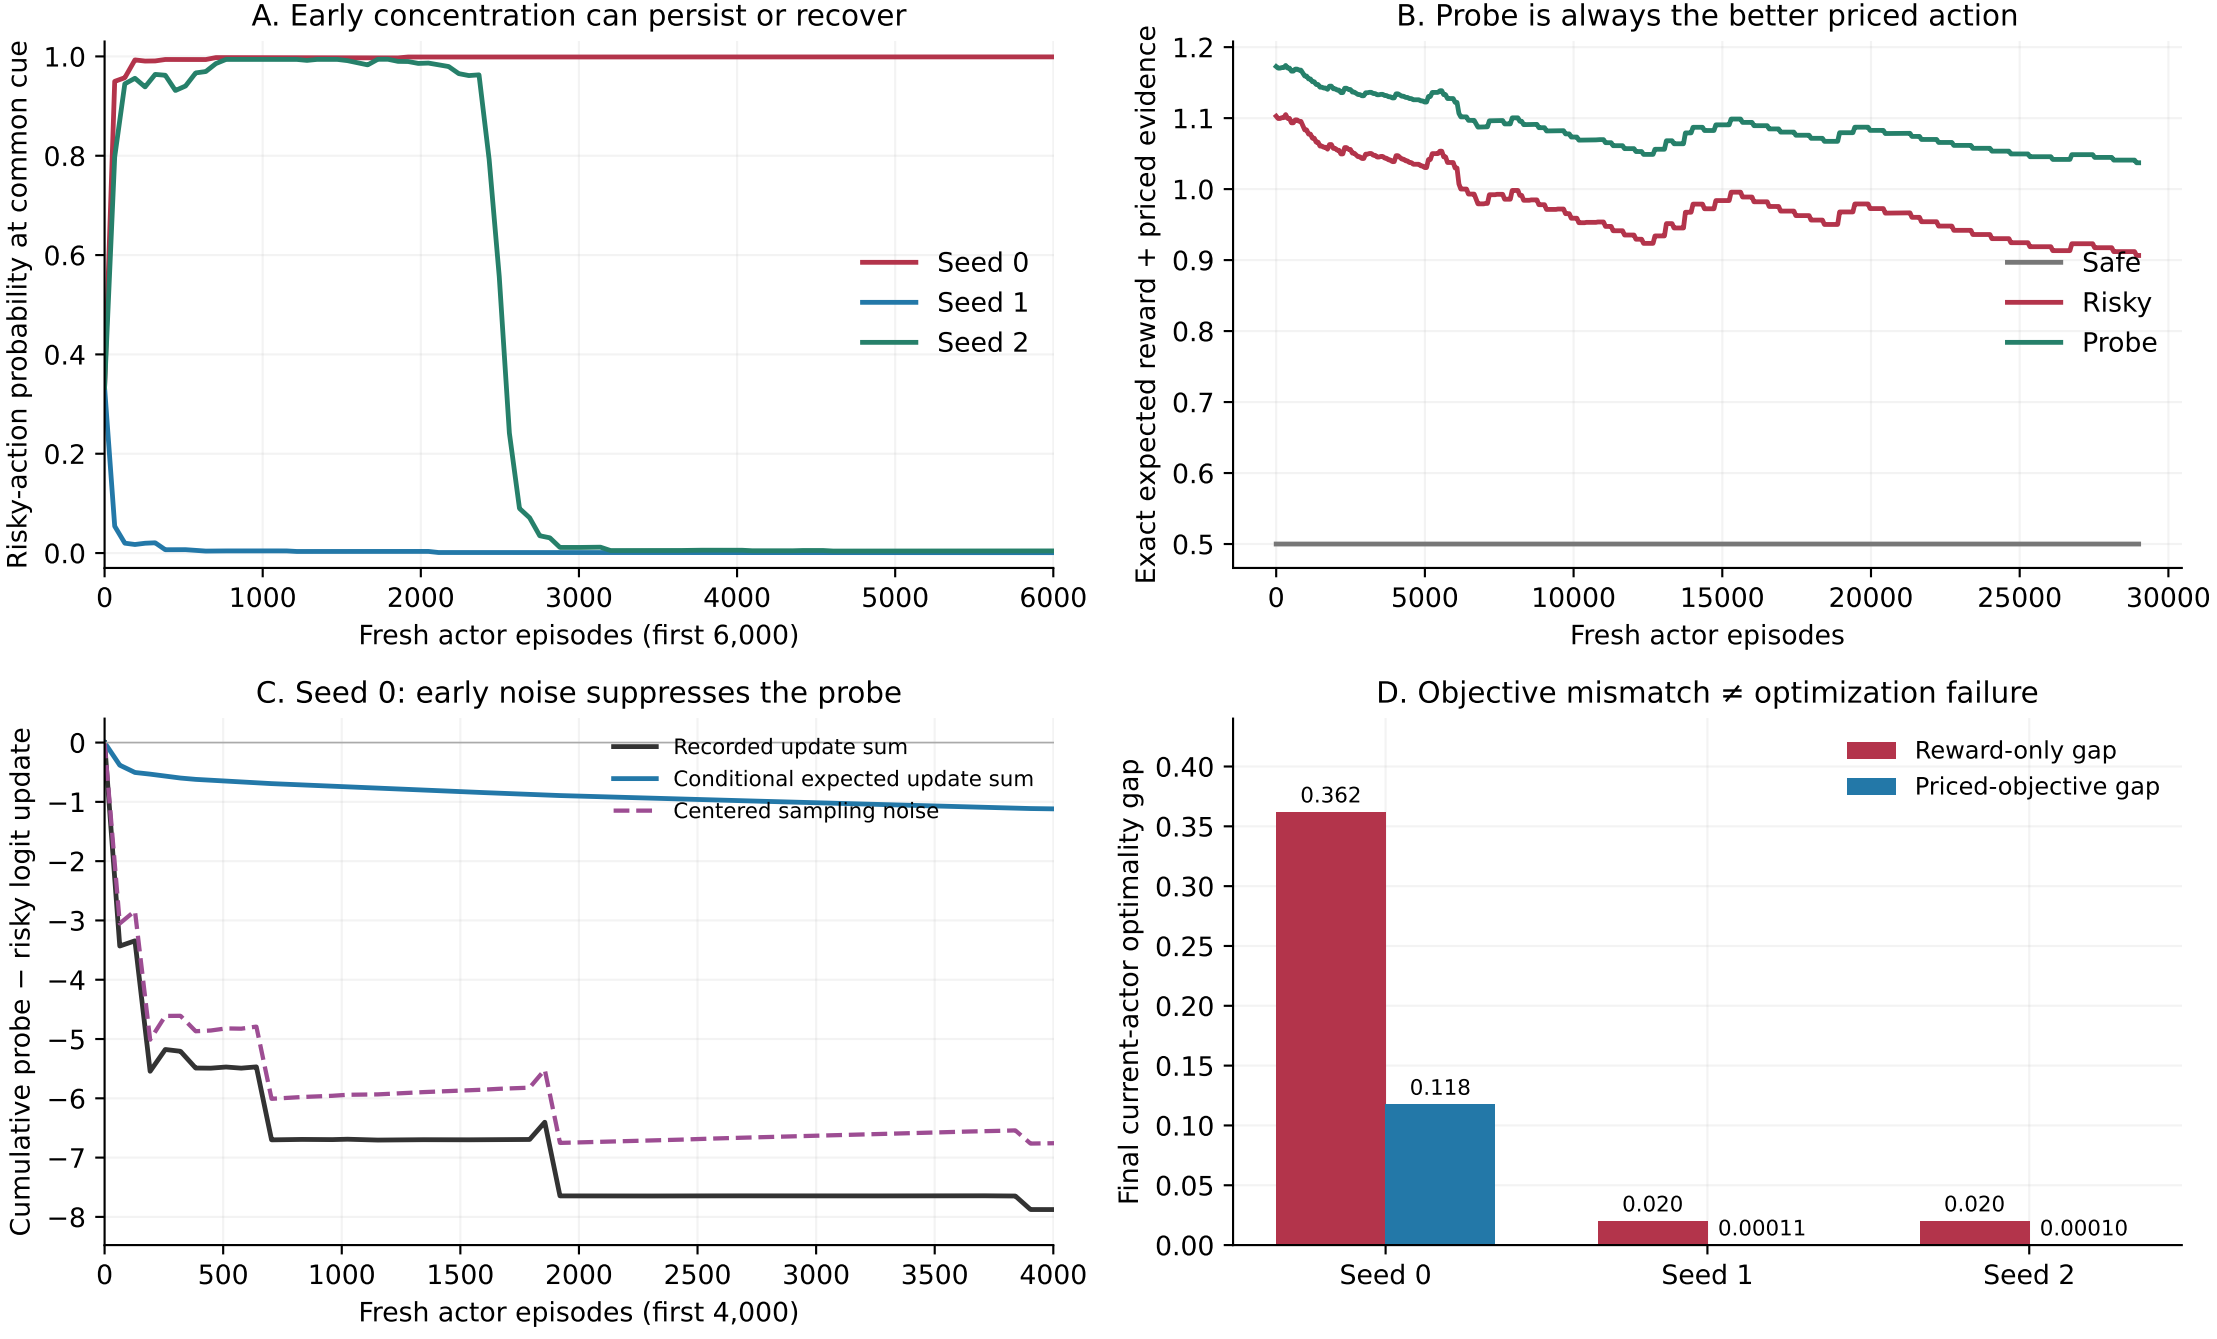
A. Early concentration can persist or recover
Risky-action probability at common cue
Fresh actor episodes (first 6,000)
Seed 0
Seed 1
Seed 2
B. Probe is always the better priced action
Exact expected reward + priced evidence
Fresh actor episodes
Safe
Risky
Probe
C. Seed 0: early noise suppresses the probe
Cumulative probe − risky logit update
Fresh actor episodes (first 4,000)
Recorded update sum
Conditional expected update sum
Centered sampling noise
D. Objective mismatch ≠ optimization failure
Final current-actor optimality gap
Reward-only gap
Priced-objective gap
0.362
0.118
0.020
0.00011
0.020
0.00010
Seed 0
Seed 1
Seed 2